

Philosophical Foundations — Module 08: Planning

— Study Questions

1. State Hume's is-ought problem. How does Active Inference relate to it?
2. What is the difference between hypothetical and categorical imperatives? Which does Active Inference provide?
3. What is the C vector's philosophical status — is it a description of what the agent expects, or a prescription of what the agent should pursue?
4. Where do preferences come from in Active Inference? Can the origin of preferences bridge the is-ought gap?
5. Compare hedonic and eudaimonic accounts of well-being. Which does Active Inference favor, and why?
6. How does the epistemic component of EFE prevent the "dark room" problem from arising in ethics?
7. Are all preferences equal within Active Inference? What criteria, if any, can distinguish "good" from "bad" C vectors?
8. How does preference coherence relate to ethical evaluation? Can incoherent preferences be considered ethically deficient?
9. How does Active Inference frame the AI alignment problem? What is the formal version of "value alignment"?
10. What is the ethical advantage of Active Inference's interpretability over black-box AI?
11. Does an AI Active Inference agent have moral status? Under what conditions?
12. How does Active Inference define responsibility? What is "model ownership"?
13. How can moral development be modeled as generative model refinement?
14. Explain guilt and shame as prediction errors within the social generative model.
15. How does Active Inference analyze political power as "precision control"?
16. Can democracy be modeled as collective Active Inference? What are the strengths and limitations of this analogy?
17. How does the concept of "justice as model alignment" differ from Rawlsian justice?
18. What does Active Inference say about moral relativism? If different cultures have different C vectors, are all moral systems equally valid?

19. How might Active Inference inform environmental ethics? Can ecosystems be modeled as collective free energy minimizing systems?
20. Critically evaluate: Does Active Inference have anything genuinely new to say about ethics, or does it merely redescribe existing ethical theories in mathematical notation?